



More than just Performance

A scale out file system that solves I/O problems

HP-CAST 24 Frankfurt

11th July 2015

Dr. Franz-Josef Pfreundt

Who makes BeeGFS ?



Fraunhofer Spin Off

Sales & Support
Industry adaptations




Research & Development


Fraunhofer Center for High Performance Computing



40 Scientists from Computer Science, Geophysics, Mathematics, Physics



Big Data
GPI-Space



Seismic Imaging
RTM – GRT – ISIM

BeeGFS Design Philosophy

- Designed for Performance and Scalability
 - Distributed Metadata
 - No Linux patches, on top of EXT, XFS, ZFS, BTRFS, ..
 - Scalable multithreaded architecture
 - Native IB and Ethernet with dynamic failover (TCP, RDMA)
 - Easy to install and maintain
 - High Software Quality
- ➔ Free software, 1000's of users, 150 supported installations

BeeGFS Key Features

■ Performance & Scalability

■ Flexibility

- Multiple daemons (any combination) can run on the same machine
- Flexible striping per file/per directory
- Add servers without downtime
- On demand filesystem „per job“ possible
- Client runs on any kernel >2.6.16
- Client runs on Xeon PHI
- ARM port available
- NFS & SMB/CIFS re-export possible



BeeGFS Key Features

- Performance & scalability
- Flexibility
- Easy to use
 - Servers run in user space
 - No kernel patches
 - Servers use existing local filesystems (ext4, xfs, zfs, ...)
 - Packages for RHEL/SL/CentOS/SLES/Debian/Ubuntu
 - Hardware independent
 - Graphical monitoring tool



Flexibility

Bursting I/O

- Checkpointing
- Open Foam
- Shared File I/O
- HDF 5
- CFD output
- Ugly Life Sciences Codes
-

Wolfgang Nagel: Flex I/O System with different I/O targets
monitoring system steers I/O

DARPA: High Speed Burst Buffer, transparent caching

Complex Solutions - Let us make it easy



BeeOND[®]

BeeGFS ON Demand


Fraunhofer Seislab

In-house cluster of CC-HPC at Fraunhofer ITWM


- 92 compute nodes with 1 TB of SSDs each
- Global BeeGFS storage on 3,5" SATA drives
- 1500 Avoton nodes , single SSD per node

How does BeeOND work

- Create a temporary PFS per Job across SSD's
- Done using PBS prolog
- Stage-in input data, work on BeeOND, stage-out results



Why ?



I/O Speed per node 100 MB/sec
Local SSD RAID 1000 MB/sec


User with 50 nodes will get 50 GB/sec



Dedicated Storage System

10 GB/sec streaming I/O

BeeOND



I/O Speed per node 100 MB/sec
Local SSD RAID 1000 MB/sec

User with 500 nodes will get 500 GB/sec


Create with BeeOND a parallel file system (BeeGFS) per
→/my_scratch/
Data Staging with beeond_cp

Dedicated Storage System



100 GB/sec streaming I/O

Shared File I/O with BeeGFS




Chunk size : 600 kB

20 server with 4x250GB SSD

Write : 52 % of maximal performance

Read : 67 % of maximal performance

IOPS low latency I/O



Mesured on a 4 year old quad SSD system

Why does this work?

- We do not patch the kernel or a underlying file system!
 - All server software is user space -> up and running in a second
 - Runs on any underlying file system (ZFS, EXT, XFS, BTTRFS,..)
 - BeeGFS server software is multithreaded and performance optimized
 - **Low CPU overhead**
 - BeeGFS scales to high server numbers : Data & Metadata
- client & server & compute on the same node



simply works

Data Availability


Data Availability Aspects of BeeGFS

- All server software is user space
 - -> Software problem would not crash a server
- Dynamic network fail over and fail back
- Uses the availability features of the underlying file systems

But

- Server Hardware breaks
- Raid Systems die
- In compute & store solutions applications may crash the system

High availability – failover with shared storage




High availability – failover with shared storage

Pros:

- No system downtime in case of server failure
- No additional storage capacity needed

Cons:

- Expensive storage components needed
- 3rd party software components needed
- Complex to set up and maintain
- Failover Risk
- Does not cover RAID
System failures




We introduce at ISC 2015

BeeGFS Enterprise Edition


BeeGFS with Built-in replication

- Assign targets to „mirror buddy groups“
- MBGs replicate chunks (but can also store non-replicated data)
- Internal HA/failover and restore mechanisms




BeeGFS with Built-in replication

- Assign targets to „mirror buddy groups“
- MBGs replicate chunks (but can also store non-replicated data)
- Internal HA/failover and restore mechanisms

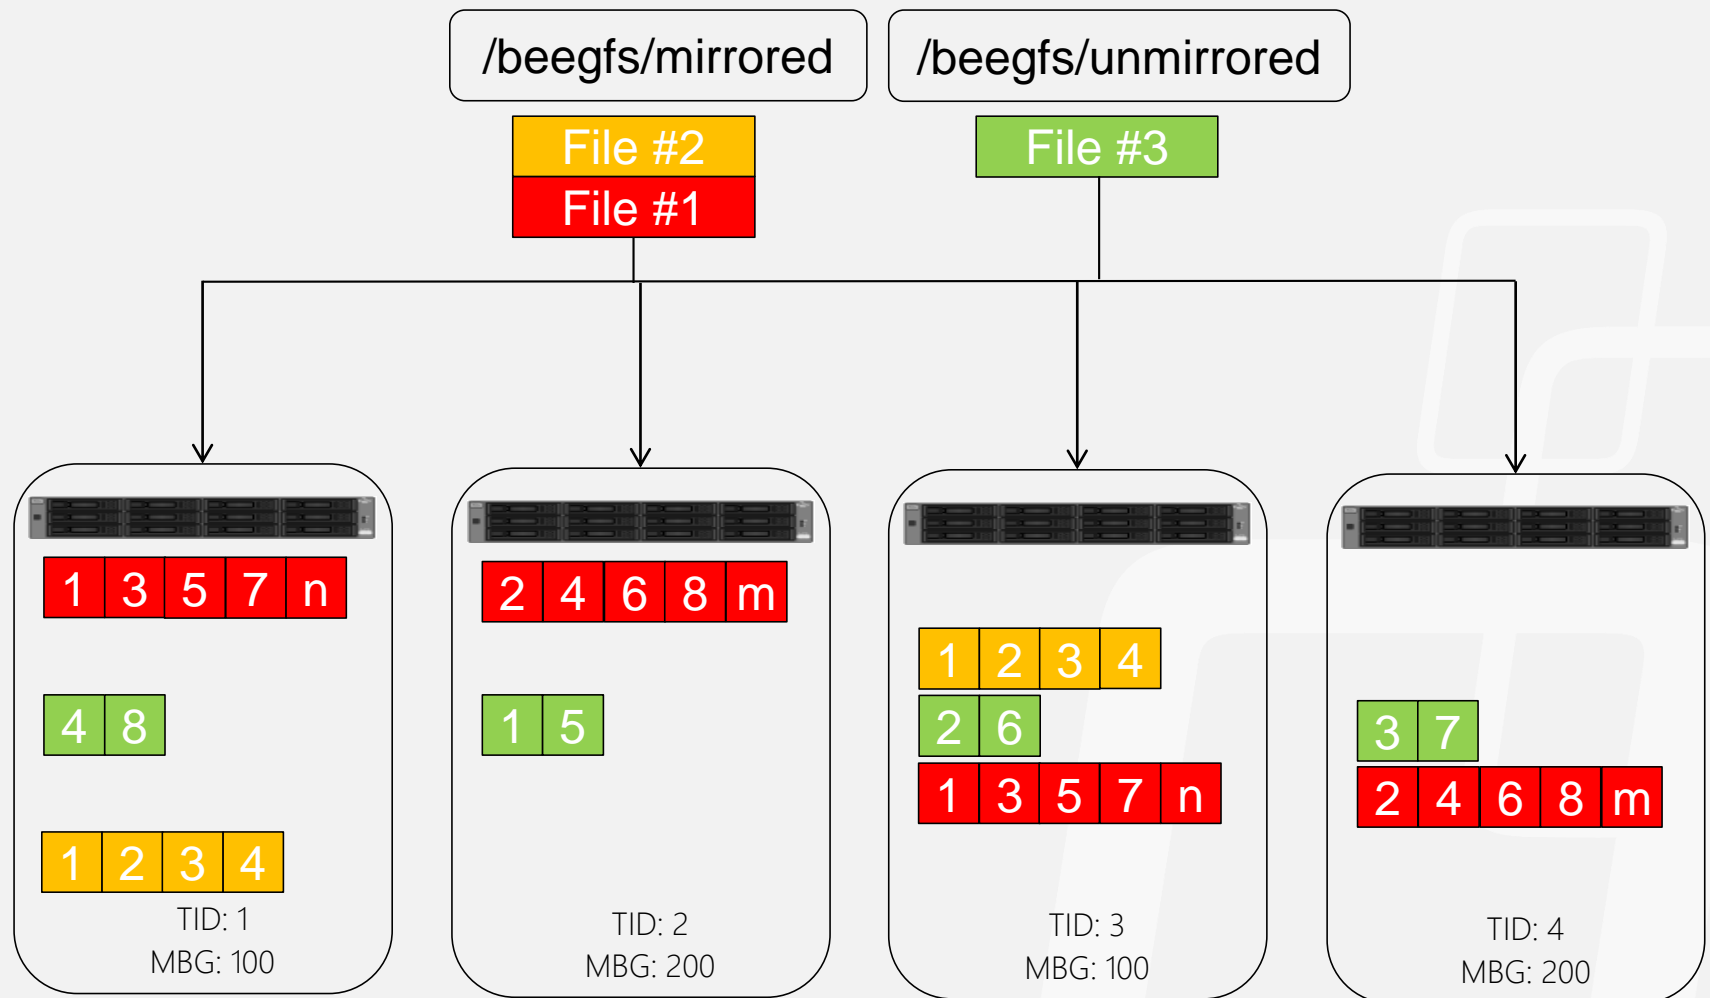


BeeGFS with Built-in replication

- Assign targets to „mirror buddy groups“
- MBGs replicate chunks (but can also store non-replicated data)
- Internal HA/failover and restore mechanisms



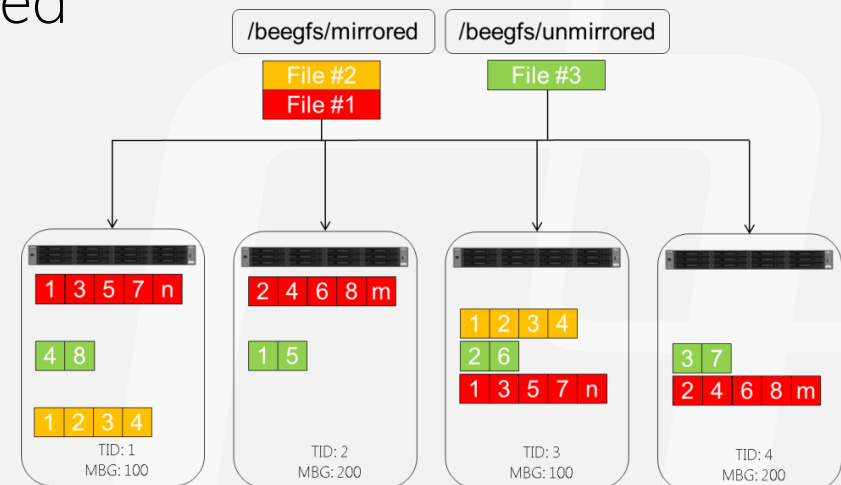
High availability – Built-in replication



High availability – Built-in replication

- Flexible (replication configurable per-directory)
- Easy to scale/extend
- No 3rd party tools for monitoring and failover functionality
- Covers server and raid failures
- Any storage backend can be used
- Additional data safety

HA at lower cost



BeeGFS EP Edition

- Built-in High Availability
- Quota
- ACL

Roadmap:


- HA for OSTs within the next weeks
- in Q4 Built-in HA for MDS

BeeGFS Hadoop Adaptor



Research Topics - BeeGFS API

Create non coherency zones using BeeOND and add a data cache



DEEP-ER

BeeGFS cache layer managed by an API (for the very few exascale apps)

BeeGFS cache API will be available in 2016

-> allows application-controlled data movements to/from global storage

Research Topics - Erasure Coding

- Add k „parity“ blocks to n stripe sets
- Reduces the amount of additional storage and covers k server failures
- $K=2$, $n=4$, 50 % more storage
- Basic library is implemented - patent free
- Performance tuning

Possible future feature?

Looking for Development Sponsors



A parallel file system that solves I/O problems

Questions?